

Lecture : SQL

Relational databases

(eg)

origin	dest	duration
NY	London	4:15
moscow	paris	4:35
:		

← each row contains data about a flight

SQL

(structured query language)

- interact with databases in an easy manner
- we'll use SQL version called

PostgreSQL

variable length characters

Data Types : integer, decimal,

(eg) string

serial, varchar,

like an integer but counts

timestamp, boolean,

automatically

enum, ...

(eg: counters)

one is a finite # of discrete values

Creating Table

files called

.sql

- create a table
called flights
rows and cols

create.sql

CREATE TABLE flights (
id	SERIAL	PRIMARY KEY
original	VARCHAR	NOT NULL
destination	VARCHAR	NOT NULL
duration	INTEGER	NOT NULL
);		

CONSTRAINTS

"CONSTRAINTS
some
properties do
data entries
in these
cols

NOT NULL - Data can't be null

UNIQUE - unique for each data point
col data

PRIMARY KEY - single way by which
table is referenced...

DEFAULT val - if I don't tell you what
to put in col, put DEFAULT
val

CHECK - only allow values that pass
CHECK conditions...

need a PostgreSQL server

start server
locally on
computer

find servers
hosted online

(e.g) Heroku

online platform
for hosting
websites and
databases

psql "url to
database"

to access database

\d

- shows different parts to database

INSERT DATA INTO TABLES

INSERT INTO

flights

table-name

(origin, destination, duration)

VALUES ('NY', 'London', 415);

names to
cols

↑
values.

Note: note that every col should be
updated...

Viewing data in a table

SELECT query

table - name
all cols
SELECT * FROM flights;

SELECT origin, destination FROM flights;

id	origin	dest	dur
=	=	=	=

restrict rows selected

SELECT * FROM flights WHERE id = 3;

more powerful queries

(eg) AVG = Average function , SUM

① SELECT AVG(duration) FROM FLIGHTS;

→ avg of all flights duration

→ SELECT AVG(duration) FROM

(2)

flights WHERE origin
= 'NewYork';

→ avg of all flights duration from NY.

(eg) COUNT → returns how many rows returned from database

SELECT COUNT(*) FROM flights;

(eg) MIN (col), MAX (col)

* can use AND, OR, IN

⋮
⋮

SELECT * FROM flights WHERE origin IN
('NewYork', 'Lima');

* string matching --- (LIKE)

SELECT * FROM flights WHERE origin

LIKE '%a%' ;



1.
 text

Updating data in a table

UPDATE flights

SET duration = 430

WHERE origin = 'New York'

AND destination = 'London';

update
duration
 $c1 = 430$
for all flights from NY → London--.

UPDATE flights

SET duration = 430

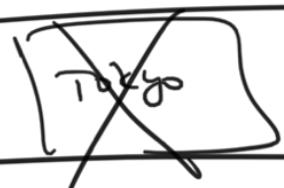
set every value in duration = 430

Delete data from a table

DELETE FROM flights

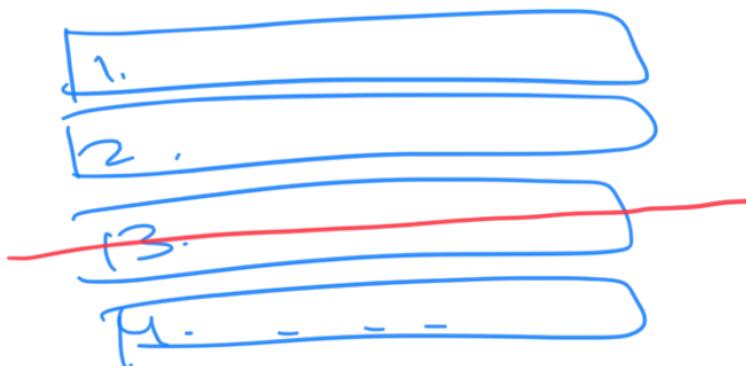
WHERE destination = 'Tokyo';

Dest

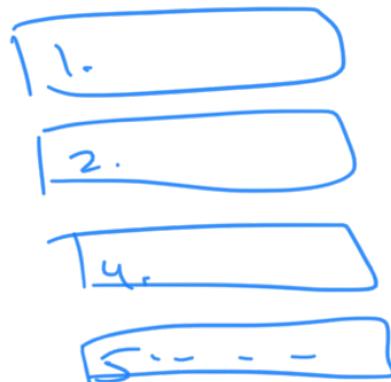


how do delete columns?

⚠ Also if we delete a row with `'id' = 3` say, no other row will ever have `'id' = 3`.



then we add a row



etc.

ADDITIONAL USEFUL COMMANDS

① Limit ← limits # of rows returned.

`SELECT * FROM flights LIMIT 2;`

② Sorting

`SELECT * FROM flights ORDER`

`BY duration ASC;`

↑
row

↑
ascending

went to
sort by

(DSC
is
descending)

Combining multiple queries into a single query

most common origins ?

SELECT origin, COUNT(*) FROM flights

GROUP BY origin

origin	count
Shanghai	1
NY	2
una	1
Istanbul	1
moscow	1

what col
you are
grouping by

SELECT origin, COUNT(*) follows FROM flights

GROUP BY origin HAVING

COUNT(*) > 1j

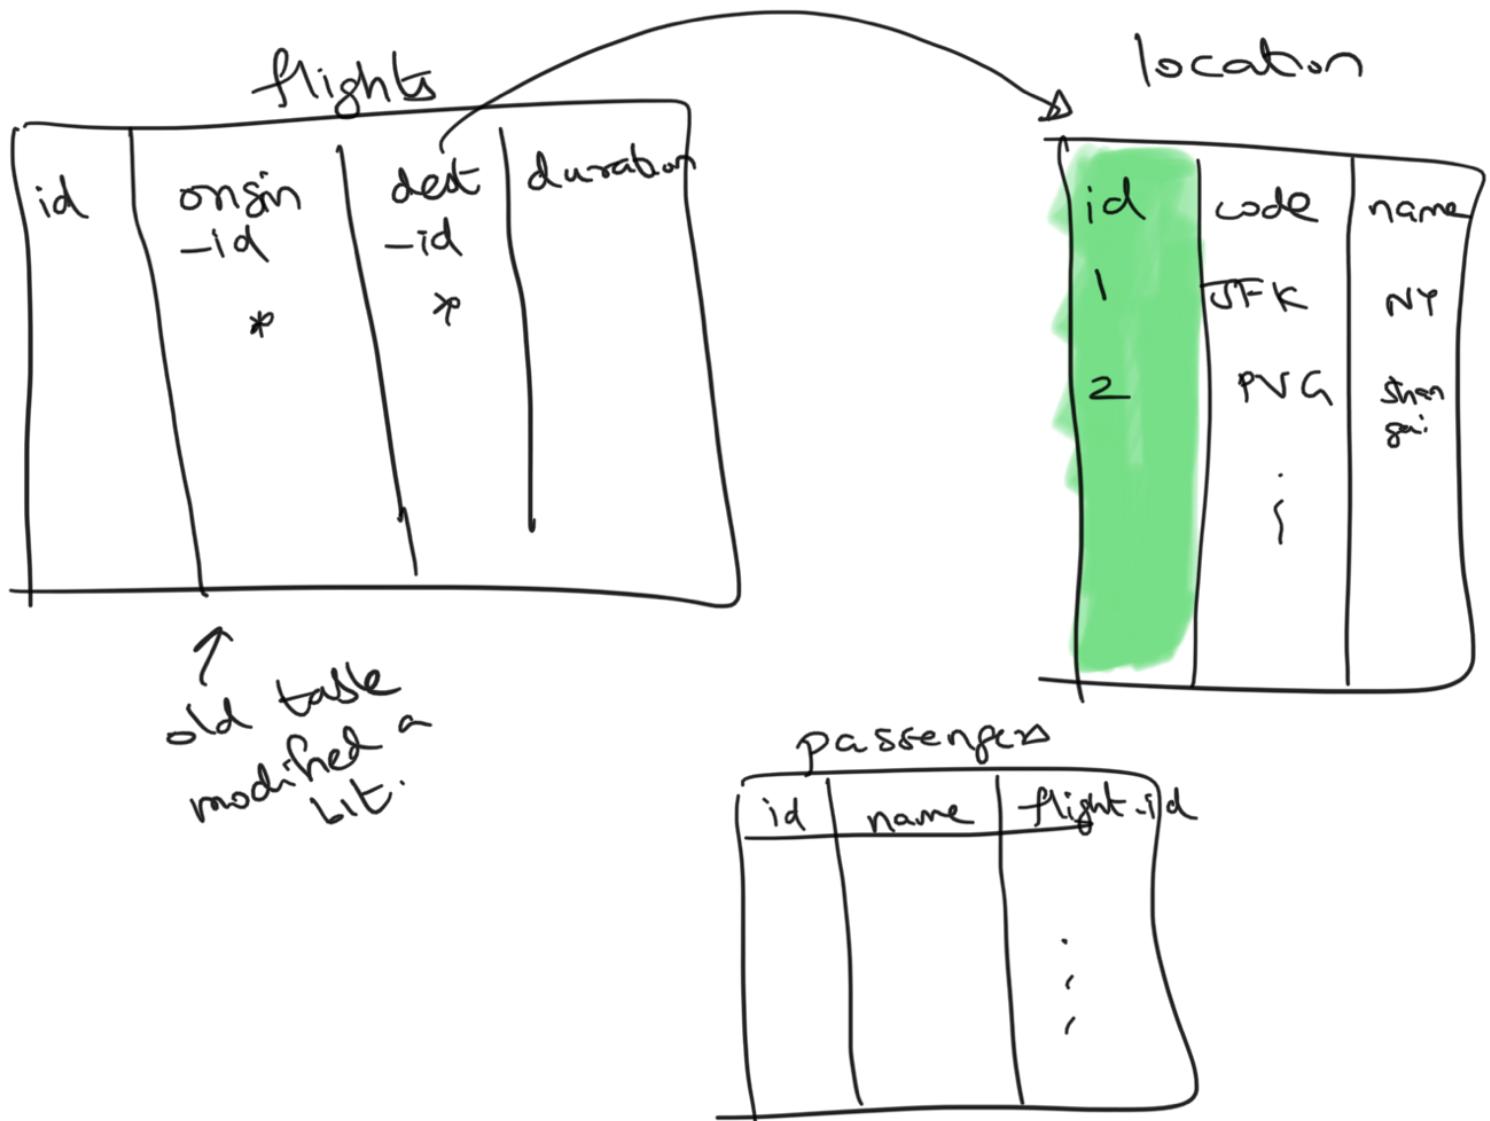
↑
what property you
want data in
origins to
tally

8u

origin	count
NY	2

Connecting multiple tables together

Foreign keys



foreign key -- "referencing key from a different table"

CREATE TABLE passengers (

etc

..

TABLE

passengers (

as many as

Implicit
referenc

passenger_id →
 id SERIAL PRIMARY KEY,
 name VARCHAR NOT NULL,
 flight_id INTEGER REFERENCES flights
 → flights

Join 2 related tables so that

you can give a more

complicated query

SELECT origin, destination, name
 → flights table passenger
 FROM flights JOIN passengers
 → which tables joining

ON passengers.flight_id

= flights.id;

relationship bet we tables

Inner Join (default)

→ returns relevant rows only if there are matches bet passengers and flights table --

flights LEFT JOIN passengers

~ matches + all rows in flights --

flights RIGHT JOIN passengers

~ matches + all rows in passengers

To include all rows in both tables ??

Creating index

- easy way to reference something by looking at something else
- can add 'index' to a particular col

(eg) if referencing flights by their origins a lot , --

index ↗ easy look up
origin | flights from origin

only

as

1, 3, 7, -

Index → takes up space
(but updating ... → takes more time)

CREATE INDEX

)
arbitrary
name-of-index

table-name

column-name

nested queries

passenger

id	name	flight-id
1	A	1
2	B	2
3	C	1
4	D	3
5	E	3
6	F	4

SELECT flight-id FROM passengers

GROUP BY flight-id HAVING

$\text{COUNT(*)} > 1$;

flight_id
1
3

SELECT * FROM flights WHERE
id IN

1 (SELECT flight_id FROM passengers
GROUP BY flight_id HAVING

$\text{COUNT(*)} > 1$;

nested query ↓ selects this -- from flights

id	origin	dest	duration
1			
2			
6			..

SQL transactions

BEGIN

— code — here

COMMIT

executed as
a bunch together
Locks database
while these are

Prevents concurrency conflicts
(RACE CONDITIONS)

SQL injection

```
SELECT * FROM USERS
WHERE (username = username)
AND (password = password);
```

+ **user-1**

' OR '1'='1

← malicious hacker
can input this

```
SELECT * FROM USERS
```

WHERE (username = '**user-1**')

AND (password = '**' OR '1'='1**');



SQL will read this as

AND **[(password = '1') OR ('1'='1')]**

which will always be True !!

① don't directly substitute user input
into SQL code directly

② sanitise input!